

# Project: Investigate a TMDb Movie Dataset

## Table of Contents

- [Introduction](#)
- [Data Wrangling](#)
- [Exploratory Data Analysis](#)
- [Conclusions](#)

## Introduction

This data set contains information about 10,000 movies collected from The Movie Database (TMDb), including user ratings and revenue.

As movies are the source of entertainment so it attract me the most how different movies earn and what factor that we should look about a movies

data have different columns :

```
['id', 'imdb_id', 'popularity', 'budget', 'revenue', 'original_title',  
 'cast', 'homepage', 'director', 'tagline', 'keywords', 'overview',  
 'runtime', 'genres', 'production_companies', 'release_date',  
 'vote_count', 'vote_average', 'release_year', 'budget_adj',  
 'revenue_adj']
```

## Question I'm intrested in:

- which movie is the highest and lowest grossing?
- popular genres year by year.
- popular genres overall
- catogorise the movies into long, medium and short
- popular actor which perform most movies.
- how profit is related by popularity, vote\_average and runtime.
- how profit rises with time.

```
In [1]: #installing wordcloud
import sys
!{sys.executable} -m pip install wordcloud
```

```
Requirement already satisfied: wordcloud in c:\users\aditya-pc\anaconda3\lib\site-packages (1.7.0)
Requirement already satisfied: numpy>=1.6.1 in c:\users\aditya-pc\anaconda3\lib\site-packages (from wordcloud) (1.18.1)
Requirement already satisfied: pillow in c:\users\aditya-pc\anaconda3\lib\site-packages (from wordcloud) (7.0.0)
Requirement already satisfied: matplotlib in c:\users\aditya-pc\anaconda3\lib\site-packages (from wordcloud) (3.1.3)
Requirement already satisfied: pyparsing!=2.0.4,!=2.1.2,!=2.1.6,>=2.0.1 in c:\users\aditya-pc\anaconda3\lib\site-packages (from matplotlib->wordcloud) (2.4.6)
Requirement already satisfied: cycler>=0.10 in c:\users\aditya-pc\anaconda3\lib\site-packages (from matplotlib->wordcloud) (0.10.0)
Requirement already satisfied: kiwisolver>=1.0.1 in c:\users\aditya-pc\anaconda3\lib\site-packages (from matplotlib->wordcloud) (1.1.0)
Requirement already satisfied: python-dateutil>=2.1 in c:\users\aditya-pc\anaconda3\lib\site-packages (from matplotlib->wordcloud) (2.8.1)
Requirement already satisfied: six in c:\users\aditya-pc\anaconda3\lib\site-packages (from cycler>=0.10->matplotlib->wordcloud) (1.14.0)
Requirement already satisfied: setuptools in c:\users\aditya-pc\anaconda3\lib\site-packages (from kiwisolver>=1.0.1->matplotlib->wordcloud) (45.2.0.post20200210)
```

The system cannot find the path specified.

```
In [2]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from wordcloud import WordCloud

from collections import defaultdict

%matplotlib inline

pd.set_option('display.max_columns', 50)

df = pd.read_csv("input/tmdb-movies.csv")
df.head(1)
```

Out[2]:

	id	imdb_id	popularity	budget	revenue	original_title	cast	homepage
0	135397	tt0369610	32.985763	150000000	1513528810	Jurassic World	Chris Pratt Bryce Dallas Howard Irrfan Khan Vi...	http://www.jurassicworld.com/

## Data Wrangling

### Step involves in wrangling

1. Remove useless columns.
2. missing value treatment.
3. Outlier removal.
4. duplicate value removal
5. change the datatypes of columns

### General Properties

```
In [3]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10866 entries, 0 to 10865
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   id                     10866 non-null  int64
1   imdb_id               10856 non-null  object
2   popularity             10866 non-null  float64
3   budget                10866 non-null  int64
4   revenue                10866 non-null  int64
5   original_title         10866 non-null  object
6   cast                  10790 non-null  object
7   homepage               2936 non-null   object
8   director               10822 non-null  object
9   tagline                 8042 non-null   object
10  keywords               9373 non-null   object
11  overview               10862 non-null  object
12  runtime                10866 non-null  int64
13  genres                 10843 non-null  object
14  production_companies   9836 non-null   object
15  release_date           10866 non-null  object
16  vote_count             10866 non-null  int64
17  vote_average           10866 non-null  float64
18  release_year           10866 non-null  int64
19  budget_adj             10866 non-null  float64
20  revenue_adj            10866 non-null  float64
dtypes: float64(4), int64(6), object(11)
memory usage: 1.7+ MB
```

```
In [4]: df.describe()
```

```
Out[4]:
```

	id	popularity	budget	revenue	runtime	vote_count	vote_average	r
<b>count</b>	10866.000000	10866.000000	1.086600e+04	1.086600e+04	10866.000000	10866.000000	10866.000000	10
<b>mean</b>	66064.177434	0.646441	1.462570e+07	3.982332e+07	102.070863	217.389748	5.974922	5
<b>std</b>	92130.136561	1.000185	3.091321e+07	1.170035e+08	31.381405	575.619058	0.935142	5
<b>min</b>	5.000000	0.000065	0.000000e+00	0.000000e+00	0.000000	10.000000	1.500000	5
<b>25%</b>	10596.250000	0.207583	0.000000e+00	0.000000e+00	90.000000	17.000000	5.400000	5
<b>50%</b>	20669.000000	0.383856	0.000000e+00	0.000000e+00	99.000000	38.000000	6.000000	5
<b>75%</b>	75610.000000	0.713817	1.500000e+07	2.400000e+07	111.000000	145.750000	6.600000	5
<b>max</b>	417859.000000	32.985763	4.250000e+08	2.781506e+09	900.000000	9767.000000	9.200000	5

```
In [5]: df.dtypes
```

```
Out[5]: id                int64
imdb_id                object
popularity            float64
budget                int64
revenue               int64
original_title        object
cast                  object
homepage              object
director              object
tagline               object
keywords              object
overview              object
runtime               int64
genres                object
production_companies  object
release_date          object
vote_count            int64
vote_average          float64
release_year          int64
budget_adj            float64
revenue_adj           float64
dtype: object
```

```
In [6]: df.columns
```

```
Out[6]: Index(['id', 'imdb_id', 'popularity', 'budget', 'revenue', 'original_title',
               'cast', 'homepage', 'director', 'tagline', 'keywords', 'overview',
               'runtime', 'genres', 'production_companies', 'release_date',
               'vote_count', 'vote_average', 'release_year', 'budget_adj',
               'revenue_adj'],
              dtype='object')
```

```
In [7]: df.shape
```

```
Out[7]: (10866, 21)
```

```
In [8]: # columns not requires in analysis
notreq = ['id','imdb_id', 'homepage', 'director', 'tagline', 'keywords', 'overview',
          'production_companies', 'vote_count', 'budget_adj', 'revenue_adj']
df.drop(notreq, axis = 1, inplace = True)
```

```
In [9]: df.columns
Out[9]: Index(['popularity', 'budget', 'revenue', 'original_title', 'cast', 'runtime',
              'genres', 'release_date', 'vote_average', 'release_year'],
              dtype='object')

In [10]: df.shape
Out[10]: (10866, 10)
```

## Data cleaning

### removing NAN values

```
In [11]: df.isnull().sum()/df.shape[0]
Out[11]: popularity      0.000000
         budget         0.000000
         revenue        0.000000
         original_title  0.000000
         cast           0.006994
         runtime        0.000000
         genres         0.002117
         release_date   0.000000
         vote_average   0.000000
         release_year   0.000000
         dtype: float64
```

```
In [12]: df.dropna(inplace = True)
```

### removing duplicate

```
In [13]: df.duplicated().sum()
Out[13]: 1

In [14]: df.drop_duplicates(inplace=True)

In [15]: df.shape
Out[15]: (10767, 10)
```

### Replace 0 with np.NAN

```
In [16]: df.describe()
```

```
Out[16]:
```

	popularity	budget	revenue	runtime	vote_average	release_year
count	10767.000000	1.076700e+04	1.076700e+04	10767.000000	10767.000000	10767.000000
mean	0.650924	1.475532e+07	4.018610e+07	102.413393	5.967549	2001.283459
std	1.003565	3.102387e+07	1.174783e+08	30.906009	0.931426	12.815909
min	0.000065	0.000000e+00	0.000000e+00	0.000000	1.500000	1960.000000
25%	0.209957	0.000000e+00	0.000000e+00	90.000000	5.400000	1995.000000
50%	0.386062	0.000000e+00	0.000000e+00	99.000000	6.000000	2006.000000
75%	0.719253	1.600000e+07	2.476490e+07	112.000000	6.600000	2011.000000
max	32.985763	4.250000e+08	2.781506e+09	900.000000	9.200000	2015.000000

```
In [17]: # these are the columns with 0 as minimum which is not possible
temp = ['budget', 'revenue', 'runtime']

df[temp] = df[temp].replace({0:np.nan})

df.dropna(inplace = True)

df.shape
```

```
Out[17]: (3850, 10)
```

### change the datatype

```
In [18]: df.dtypes
```

```
Out[18]: popularity      float64
budget      float64
revenue      float64
original_title  object
cast         object
runtime      float64
genres       object
release_date  object
vote_average  float64
release_year  int64
dtype: object
```

```
In [19]: # release date: object -> date
df['release_date'] = pd.to_datetime(df['release_date'])
```

```
In [20]: # change : float to int
change = ['popularity', 'budget', 'revenue']

df[change] = df[change].astype('int')

df.dtypes
```


```
Out[20]: popularity      int32
budget      int32
revenue      int32
original_title  object
cast      object
runtime      float64
genres      object
release_date  datetime64[ns]
vote_average  float64
release_year  int64
dtype: object
```

```
In [21]: f"There are total {df.shape[0]} movies"
```


```
Out[21]: 'There are total 3850 movies'
```

## Exploratory Data Analysis

```
In [22]: df.hist(figsize=(8,8));
```



```
In [23]: # i have used heatmap to see the correlation
corr = df.corr()
plt.figure(figsize=(8,8))
ax = sns.heatmap(
    corr,
    vmin=-1, vmax=1, center=0,
    cmap=sns.diverging_palette(20, 220, n=200),
    square=True
)
ax.set_xticklabels(
    ax.get_xticklabels(),
    rotation=45,
    horizontalalignment='right'
);
```



## Q1 : which movie is the highest and lowest grossing?

```
In [24]: df['profit'] = df['revenue']-df['budget']
# used this to calculate profits
```

```
In [25]: highest_profit = df['profit'].idxmax()
lowest_profit = df['profit'].idxmin()
#used to find the index of maximum and minimum
```



```
In [26]: ext = pd.DataFrame([df.iloc[highest_profit,:], df.iloc[lowest_profit,:]])
ext
```

Out [26]:

	popularity	budget	revenue	original_title	cast	runtime	genres	release_year
4021	0	50000000	26199517	Blood Work	Clint Eastwood Jeff Daniels Anjelica Huston Wa...	110.0	Crime Drama Mystery Thriller	2002
6559	3	12000000	114194847	Step Up	Channing Tatum Jenna Dewan Damaine Radcliff De...	104.0	Music Drama Romance Crime	2006

## popular genres of that year

```
In [27]: df1 = df.copy(deep=True)
```


```
In [28]: df1['genres'] = df1['genres'].apply(lambda x: " ".join(str(i)+" " for i in x.split('|')))
```

```
In [29]: years = df1.groupby('release_year')['genres'].sum()
```

```
In [30]: def cloud(year):
    wordcloud = WordCloud(width = 800, height = 800,
                           background_color = 'white',
                           min_font_size = 10).generate(years[year])
    plt.figure(figsize = (6, 6), facecolor = None)
    plt.imshow(wordcloud)
    plt.axis("off")
    plt.tight_layout(pad = 0)
    plt.show()
```

```
In [31]: year = int(input("Enter year in b/w 1960-2015 "))  
cloud(year)
```

Enter year in b/w 1960-2015 2015



popular genres overall


```
In [32]: # create the frequency dictionary
gen = df['genres'].str.cat(sep='|')

gen = pd.Series(gen.split("|"))

a = gen.value_counts(ascending=False)
a
```

```
Out[32]: Drama          1754
Comedy          1358
Thriller        1203
Action          1085
Adventure        749
Romance          667
Crime            651
Science Fiction  519
Horror           463
Family           425
Fantasy          396
Mystery          344
Animation        201
Music            134
History          129
War              119
Western           52
Documentary       31
Foreign           13
TV Movie          1
dtype: int64
```

```
In [33]: wordcloud = WordCloud(width = 800, height = 800,  
                                background_color = 'white',  
                                min_font_size = 10).generate_from_frequencies(a)  
plt.figure(figsize = (6, 6), facecolor = None)  
plt.imshow(wordcloud)  
plt.axis("off")  
plt.tight_layout(pad = 0)  
plt.show()
```



**categorise the movies into Extra long, long, medium and short**

```
In [34]: bins = [15, 95.25, 106, 119, 338]
```

```
In [35]: labels = ['Extra long', 'Long', 'Medium', 'short']
```


```
In [36]: df['category'] = pd.cut(df['runtime'],bins,labels=labels)
df.head()
```

Out [36]:

	popularity	budget	revenue	original_title	cast	runtime	genres	release_
0	32	150000000	1513528810	Jurassic World	Chris Pratt Bryce Dallas Howard Irrfan Khan Vi...	124.0	Action Adventure Science Fiction Thriller	2015-0
1	28	150000000	378436354	Mad Max: Fury Road	Tom Hardy Charlize Theron Hugh Keays-Byrne Nic...	120.0	Action Adventure Science Fiction Thriller	2015-0
2	13	110000000	295238201	Insurgent	Shailene Woodley Theo James Kate Winslet Ansel...	119.0	Adventure Science Fiction Thriller	2015-0
3	11	200000000	2068178225	Star Wars: The Force Awakens	Harrison Ford Mark Hamill Carrie Fisher Adam D...	136.0	Action Adventure Science Fiction Fantasy	2015-1
4	9	190000000	1506249360	Furious 7	Vin Diesel Paul Walker Jason Statham Michelle ...	137.0	Action Crime Thriller	2015-0

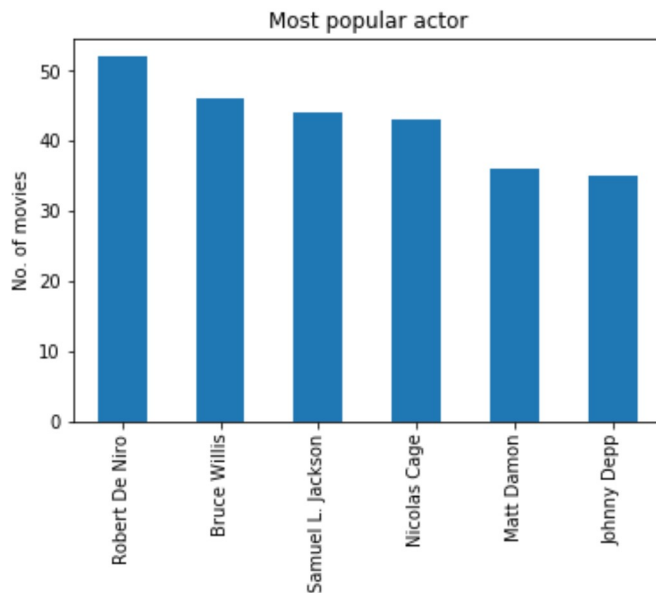
```
In [37]: df['category'].value_counts().plot(kind='bar')
plt.ylabel('runtime')
plt.title("Quantity of various category")
```

Out [37]: Text(0.5, 1.0, 'Quantity of various category')




popular actor which perform most movies.

```
In [38]: x=df['cast'].str.cat(sep="|")
cast = pd.Series(x.split("|"))
cast.value_counts(ascending=False)[:6].plot(kind='bar')
plt.ylabel("No. of movies")
plt.title('Most popular actor');
```

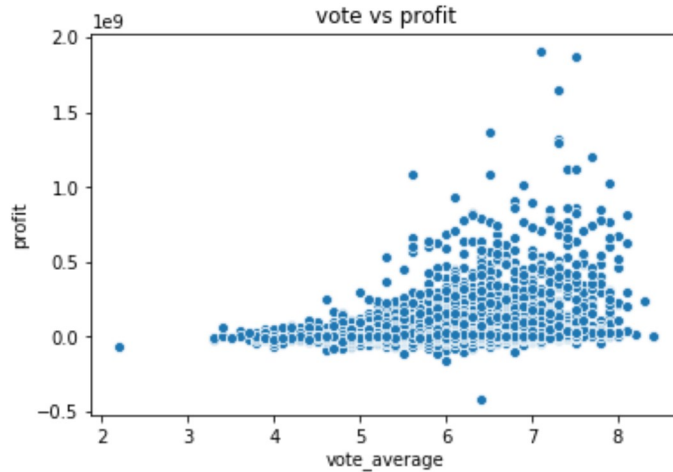


how profit is related by popularity, vote\_average and runtime.

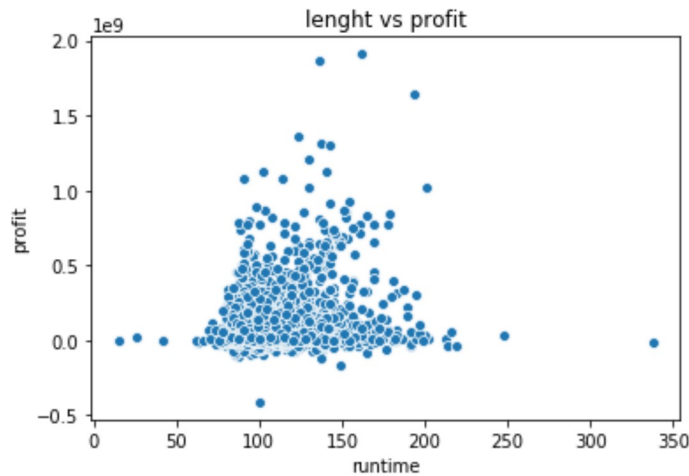
```
In [39]: sns.scatterplot(x=df['popularity'],y=df['profit'])
plt.title("popularity vs profit");
```



```
In [40]: sns.scatterplot(x=df['vote_average'],y=df['profit'])  
plt.title("vote vs profit");
```




```
In [41]: sns.scatterplot(x=df['runtime'],y=df['profit'])  
plt.title("lenght vs profit");
```



**how profit rises with time.**

```
In [42]: year = df.groupby('release_year')['profit'].sum()
plt.figure(figsize=(10,5))
plt.plot(year)
plt.title('Profit sorring as high with time')
plt.xlabel("year")
plt.ylabel("profits");
```



## Conclusions

- maximum profit by setup and maximum loss by bloodwork
- in 2015 maximum movies are of drama
- Drama is the most popular genres overall
- we have most of the movies as long followed by extra long
- Robert de nitro made most of the movies
- profits increases as year pass by

```
In [43]: from subprocess import call
call(['python', '-m', 'nbconvert', 'Investigate_a_Dataset.ipynb'])
```

Out[43]: 0